

SEMANTIC WEB PROJECT

VN-Edu LOD

A Knowledge Graph of Vietnamese Higher-Education Institutions

Design, Implementation, and Critical Evaluation

Revised English technical report

Ontology 2.2 and integration edition: 8 October 2026

This edition incorporates an independent code and data audit of the repository and recomputes measurements from the corrected release rather than copying them from the earlier Vietnamese report. The earlier report remains an archival description; its accuracy, security, and performance claims are not evidence for this edition.

Project repository:

<https://github.com/pham-ng/Vietnam-University-Knowledge-Graph-ver2>

Abstract

VN-Edu LOD integrates information about Vietnamese educational institutions into an RDF knowledge graph, combining structured Wikidata records, Vietnamese Wikipedia infoboxes, and curated reference tables. The corrected snapshot contains 300 educational-organization records, of which 271 are classified as higher-education institutions under the project's classification rules. It also represents 1,702 person records, 34 current and 29 former province-level units, and 78 illustrative academic programs. An ontology supports organizational, geographic, and educational relations; rule-based materialization adds inverse relations, classifications, and geographic containment. JSON Schema and SHACL enforce selected data contracts.

The resulting serving graph contains 73,628 distinct triples. An independent review identified defects in SPARQL request validation, query cancellation, publication gates, entity identity, persistent URI allocation, and the interpretation of evaluation results. The revised implementation addresses these defects and adds targeted regression tests. Evaluation distinguishes local structural validity, reference-link agreement, and factual accuracy. The ontology passes the OWL API OWL 2 RL profile gate with zero violations. A dated seven-request publication check has limited scope. The result is a reproducible educational Linked Data system with explicit operational and epistemic limits.

Keywords: knowledge graph; RDF; ontology; SHACL; SPARQL; entity resolution; data provenance; Vietnamese higher education.

Contents

1	Introduction and scope	5
2	Technical foundations	5
3	Data sources and collection	6
3.1	Source roles	6
3.2	Scope and selection effects	6
4	Architecture and data management	6
4.1	Layers and contracts	6
4.2	Identity and URI persistence	7
4.3	Provenance and disputed values	8
5	Transformation: from source records to RDF	8
5.1	Integration policies and information loss	8
5.2	Concrete mapping example	8
5.3	Linking, reasoning and release boundaries	10
6	Ontology design and relationship semantics	10
6.1	Modules	10
6.2	Visual specification of ontology 2.2	11
6.3	Implemented corrections and evidence limits	16
6.4	Standards and conformance assessment	16
6.5	Qualified observations and temporal interpretation	17
6.6	Competency questions and prohibited conclusions	18
6.7	Materialization	18
6.8	Geographic and temporal limitations	19
7	Linking and publication	19
7.1	External relations	19
7.2	Representations and deployment	19
8	URI publication and Apache Jena Fuseki	20
8.1	Identity, representation and storage	20
8.2	Safe loading and reproducible runtime	20
8.3	Production-like security and performance acceptance	21
8.4	Executed end-to-end acceptance evidence	21
9	Five-star Linked Open Data assessment	22

9.1 Dated HTTP evidence	23
10 Query execution and security review	23
11 Evaluation and results	24
11.1 Release measurements	24
11.2 Link experiment: agreement, not independent accuracy	24
11.3 Executed Silk Framework experiment	25
11.4 Regression evidence	26
11.5 Treatment of earlier comparisons	26
12 Threats to validity and remaining work	26
13 Conclusion	27
A Reproducing the release and report	27
B A competency-query example	29

List of Figures

1	Validated release flow	7
2	Transformation stages and data artifacts	9
3	Protégé OntoGraf view of ontology 2.2	11
4	GraphDB class-hierarchy view	12
5	GraphDB visual graph for University	12
6	Protégé object-property signature	13
7	GraphDB domain–range relation graph	14
8	Organizational class hierarchy	14
9	Property signatures across ontology modules	15
10	Geographic property-chain inference	15
11	Evidence-qualified temporal observations	18
12	URI publication and optional Fuseki execution	21

1 Introduction and scope

Information about a university can appear under different names and identifiers across encyclopedic articles, knowledge bases, and organizational websites. Relationships are equally important: a member school is different from a branch campus; direct governance is different from general state oversight; and the area containing an institution may have changed its administrative representation. A knowledge graph makes these distinctions explicit and supports queries that join several kinds of entity.

This project is an engineering case study of that integration problem. Its contribution is the combination of a domain model, a layered data pipeline, machine-readable outputs, and a query interface. It does not introduce a new reasoning algorithm or establish a complete official register of Vietnamese higher education. Names, classifications, and dates reflect the collected snapshot and documented extraction policies.

The evaluation is organized around four questions:

1. Can the model distinguish relevant organizational relationships and support useful multi-step queries without introducing contradictions detected by the implemented checks?
2. Can the pipeline reproduce a validated release and preserve stable identifiers as labels and input ordering change?
3. What do the available data and evaluation procedures establish about completeness, link agreement, and correctness, and what remains unmeasured?
4. Can publication and public query handling fail safely on invalid data and hostile inputs?

The audit covers source code, the ontology, SHACL shapes, cached inputs, generated RDF, tests, and the existing report. Security counterexamples are executed locally; external production endpoints are not used as attack targets. The project author's identity, institution, course identifier, and supervisor are not inferred or fabricated.

2 Technical foundations

RDF describes resources through triples. Subjects may be IRIs or blank nodes; predicates are IRIs, and objects may additionally be literals. A graph is a set, so duplicate triples do not increase its size. This matters when combining asserted data, ontology statements, links, and inferred statements: their file counts cannot simply be added to obtain the size of the union [1].

OWL axioms express implications about the modeled world. A property's domain is not an input-validation rule: using the property can entail that its subject belongs to that domain. Similarly, a functional object property can imply equality between two values; it does not necessarily reject the second value. The project therefore supplements reasoning with explicit validation of asserted records. The OWL 2 profile specification defines restrictions under which particular reasoning guarantees apply [2].

SHACL validates a supplied data graph against declared constraints [3]. Conformance is relative to those constraints, not a claim that every statement is true. Here, missing values for some descriptive fields generate warnings, while invalid datatypes, selected cardinality violations, and hard constraints prevent publication.

SPARQL supplies query syntax and algebra for RDF, including optional patterns, aggregates, and remote SERVICE clauses [4]. Exposing it over HTTP introduces an execution boundary: query syntax, remote destinations, elapsed work, and publication state must all be controlled independently of the user interface.

The project reuses identifiers and vocabulary from Wikidata, DBpedia, schema.org, FOAF, SKOS, and PROV-O. PROV-O supports descriptions of derivation [5]; whether the provenance is attached to an entity, a graph, or an individual assertion determines what can actually be traced. HTTP identifiers and useful representations support Linked Data publication, while persistence requires more than generating readable slugs [6].

3 Data sources and collection

3.1 Source roles

Wikidata contributes identifiers, labels, institutional attributes, organizational links, and person-related records. Vietnamese Wikipedia contributes infobox fields, descriptions, historical text, and links between articles. DBpedia supplies candidate identity links and some fallback values. Selected reference tables describe province mergers, regional groupings, fields of study, majors, and example programs. Geocoding is used as a fallback for location resolution and is not equivalent to an authoritative campus register.

Wikipedia and Wikidata are interconnected community resources. Agreement between them is useful corroborating evidence, but they must not be described as statistically independent sources. Imported statements can share an origin, and one resource may repeat an error from another. A disagreement log exposes selected conflicts; it does not establish which source is correct.

The collection strategy combines article-category traversal, infobox parsing, and Wikidata queries. Results are retained in JSON snapshots. The repository also ships a compressed HTTP cache. In the revised default pipeline, a cache miss raises an error instead of silently contacting a changing remote source. A fresh collection is an explicit operation. Dependency installation may still require Internet access; the offline claim applies to cached data collection and reconstruction after dependencies are installed.

3.2 Scope and selection effects

The snapshot is a discovered set of educational organizations, not a census. Filtering removes unrelated articles and some records lacking sufficient cross-source evidence. That policy can exclude legitimate but poorly documented institutions. Institutions with more detailed encyclopedia coverage are more likely to have richer records. The same bias is particularly strong for people linked through education records.

The retained higher-education classification includes name-based heuristics. A title beginning with a familiar institution type is evidence used by the implementation; it is not a substitute for a current legal decision. Counts of higher-education institutions must therefore be read as counts under the project's classification policy.

Academic programs are illustrative hand-entered data. They demonstrate program-to-major queries but do not support claims about a university's complete or current admissions offerings. Similarly, a Wikidata education relation establishes recorded attendance or education, not necessarily completion of a degree. The existing *Alumnus* class must be interpreted with that limitation.

4 Architecture and data management

4.1 Layers and contracts

Bronze stores responses and extraction snapshots. Silver contains institutions, people, governing bodies, and provinces in a normalized JSON representation. Gold contains asserted RDF, external links,




Figure 1: Release flow. A failed validation prevents replacement of validated inference and the serving graph; changed inputs invalidate the publication fingerprint.

selected instance-level entailments, and dataset metadata. The layers improve inspectability, but the implementation rewrites snapshot files during reconstruction. Bronze is therefore not an immutable append-only archive; Git revisions provide versioned snapshots.

Silver validation checks schema constraints and selected cross-field rules. The audit adds checks for future years, impossible calendar dates, missing cross-collection references, and invalid leadership references. These checks complement, rather than replace, semantic review of the extraction policy.

After RDF generation, literal ranges and functional-field cardinalities are validated on the raw asserted graph, independently of equality in the closure. The rule engine’s own error messages are retained in the consistency diagnostics. SHACL then checks the asserted and retained inferred graph. A failure stops the release step. On success, a validation record binds the checked release to its ontology, shapes, Silver files, component RDF files, and serving graph. The publisher rejects missing or stale validation records.

Fingerprints normalize text line endings to LF so that Git checkout on Windows and Linux does not create false stale-release failures. This is a content-integrity check, not a cryptographic signature or a guarantee against someone deliberately editing both the inputs and validation record. Reproduction of RDF is evaluated semantically; timestamps, blank-node labels, and serialization order can change without changing graph meaning.

4.2 Identity and URI persistence

A readable slug is used when a resource is first assigned a local URI. The revised minter consults the committed `uri_map.json` before allocating anything new. Renaming an entity or inserting another entity with the same label does not change a registered URI. Retired identifiers remain reserved. Source-key changes, such as a new QID replacing an article-based key, still require an explicit identity migration policy.

The previous person integration merged unidentified leaders by normalized personal name across all institutions. The audit found six names used across multiple institutions; splitting these records adds

seven person records. This does not prove that seven new real people exist. It represents uncertainty conservatively: records without a shared identifier are separated by institution until identity can be established. The former ambiguous URI is reserved but is not asserted as identical to either replacement.

4.3 Provenance and disputed values

Institution records use `prov:wasDerivedFrom` to reference Wikidata items and specific Wikipedia revisions. Conflict and fallback logs provide additional evidence for some chosen fields. This is entity-level provenance plus partial field-level logging; the project does not yet attach complete provenance to every triple. A Wikidata entity URI is also not a pinned statement revision.

The earliest reported foundation year is the current integration policy and may refer to a predecessor. Five records contain a different year in the infobox foundation date. The revised exporter retains such a date under `vnedu:reportedFoundingDate`, rather than asserting it as an unqualified `schema:foundingDate`. This preserves evidence without pretending that the historical ambiguity has been adjudicated. Distinguishing predecessor origin, legal establishment, and opening would require dated event records and authoritative supporting sources.

5 Transformation: from source records to RDF

Transformation is implemented in Python, not in Fuseki. The main pipeline calls `step3_integrate.py` to reconcile source records and `step3_transform.py` to construct RDF with `RDFLib`. It does not use RML, R2RML, an Apache Jena mapping engine, or a relational database. Figure 2 connects the operations to their actual artifacts.

5.1 Integration policies and information loss

The integration script joins articles through Wikidata identifiers where available, normalizes names, parses dates and quantities, and resolves organization references. When a target is an educational institution, it becomes a membership or branch relation; a recognized public body becomes a governance target; and a company becomes an ownership target. Unknown targets go to an unresolved-reference report. A ministry associated with a record classified as private is moved to `stateManagedBy`. These are executable heuristics, not legal determinations. The original source predicate and temporal scope are not preserved on every resolved edge.

The earliest candidate founding year is selected and disagreements are logged. Integer counts lose the measurement context when no corresponding year or unit is retained. Some campus and funding values stay as source strings because their units are not reliably normalized. The source collector's loss of birth-date precision cannot be repaired simply by assigning an RDF datatype later. Bronze, Silver and conflict reports must therefore be retained together when reviewing a disputed result.

5.2 Concrete mapping example

Table 1 follows the shipped Silver record `Q3075696`, Hanoi University of Science and Technology. The table demonstrates transformation, not an independent verification of the institution's current legal status or leadership. Let `u:` abbreviate the project's `resource/university/` namespace.

The exporter also emits coordinates as decimals, counts as non-negative integers, and websites as IRIs. Missing optional values usually produce no triple; they are not encoded as an empty string, zero, or a claim of nonexistence. Reference programs generate a program URI from the registered institution slug and major code. Consequently, this model identifies one illustrative offering per institution-major




Figure 2: Implemented transformation and validation stages. Dashed arrows denote auxiliary inputs, not an independent source of factual verification. Fuseki is an optional consumer of the final graph, not a transformation step.

Table 1: Selected Silver-to-RDF mappings for a real input record.

Silver input	RDF output	Rule or caveat
key = Q3075696	Subject u:dai-hoc-bach-khoa-ha-noi	Reuse the persistent registry; do not mint again on a label change.
kind = University	rdf:type vnedu:University	Classification came from integration policy, not from Fuseki.
name_vi	rdfs:label and skos:prefLabel, tagged @vi	Keep language tags, not an untyped display string.
founding_year = 1956	vnedu:foundingYear "1956" ^^xsd:integer	May refer to a predecessor's origin; integer is used to remain inside OWL 2 RL.
province = Q1858	vnedu:locatedIn the registered Hanoi URI	An object IRI, not a province-name literal.
ownership = public	vnedu:ownership vnedu:PublicOwnership	Later classifies the subject as a public institution.
viwiki_revid = 75604914	prov:wasDerivedFrom a Wikipedia oldid URL	Revision-specific entity provenance; not per-triple attribution.

pair; it cannot distinguish multiple curricula, degree levels, delivery modes or academic years for that pair without extending the identifier design.

5.3 Linking, reasoning and release boundaries

`step4_link.py` writes a separate link graph and dataset metadata. The reasoner reads the local ontology plus asserted data, computes a closure, and retains selected instance-level entailments. Raw asserted constraints are evaluated against the unmerged asserted graph, independently of equality effects. The current script invokes this check after closure computation, but before replacing the inferred and serving files; it is not an early pre-computation performance gate. Successful validation produces `validated-release.json`. Publication checks the recorded input hashes.

The serving file is the set union of ontology, assertions, retained inferences, links and metadata. Separate files remain available, but the union itself does not carry named-graph boundaries. Loading that file into Fuseki’s default graph does not restore those boundaries or create assertion-level provenance. A versioned dataset with named graphs and an explicit default-union query policy would be a future design change, not an existing feature.

This is a staged batch workflow, not a transaction spanning every artifact. Failed validation preserves the previous inferred and serving files, while preceding stages may already have rewritten candidate assertions and links. Hash checks protect the static publisher; the revised Fuseki loader enforces that same release gate. An immutable staging directory followed by an atomic release-pointer switch would make this boundary stronger.

6 Ontology design and relationship semantics

6.1 Modules

The ontology has 39 named local classes, 38 local object properties, and 27 local datatype properties in the audited release. Table 2 summarizes the main relationship choices. Figures 8–10 visualize selected implemented axioms, not a proposed replacement ontology. Classes, property signatures and individual-level inferences are distinguished explicitly.

Table 2: Main modules and intended distinctions.

Module	Main entities or properties	Interpretation
Organizations	Institution, governing body, company; memberOf, branchOf, governedBy, ownedBy	Membership, branch status, governance, and ownership are separate relations.
Oversight	stateManagedBy	General state oversight is distinct from direct governance.
People	Person, InstitutionLeader, Alumnus; rector, director, councilChair	Leadership roles differ; education records do not certify graduation.
Geography	Province, FormerProvince, Region; locatedIn, partOf, mergedInto	Geographic crosswalks support aggregation under successor boundaries.
Education	AcademicProgram, Major, FieldOfStudy; offersProgram, ofMajor, inField	Example programs connect institutions to subjects; these are not an official offering inventory.
Identity	owl:sameAs, skos:closeMatch, foaf:isPrimaryTopicOf	Identity, approximate concept mapping, and document-topic links are not interchangeable.

Disjointness axioms separate selected organizational and geographic categories. Their usefulness depends on defensible category definitions. A statement that violates a disjointness axiom could indicate erroneous data or an overly rigid ontology, so a detected conflict should trigger inspection rather than automatic deletion of evidence.

Vocabulary alignments use subclass and subproperty relationships. Such alignments still carry logical implications; choosing a weaker alignment than equivalence does not make it semantically harmless. The current inference run does not import and validate the complete external ontologies. Its absence of detected conflicts therefore applies to the local reasoning scope, not every possible external combination.

6.2 Visual specification of ontology 2.2




Figure 3: Direct screenshot of Protégé 5.6.7 with the OntoGraf plug-in, after loading the committed ontology/vnedu.ttl. The class hierarchy is shown at left and the selected ontology graph at right; the graph was laid out by OntoGraf and includes the selected OWL classes and their asserted subclass relations.

The two relation views complement one another. Protégé exposes the OWL property signature and makes the domain, range, inverse and functional characteristics auditable one property at a time; GraphDB renders the same class–property–class pattern as a navigable network. OntoGraf’s current plug-in session rendered the subclass hierarchy reliably but did not expose all domain–range arcs in its Arc Types panel, so this report does not mislabel the GraphDB network as a Protégé screenshot.




Figure 4: Direct screenshot of GraphDB Workbench 11.1.0 after loading the committed ontology into the local repository vnedu-ontology-d. The class-hierarchy view reports 57 classes and displays the hierarchy and reused vocabulary classes. This is a repository-backed visualization, not a diagram generated outside the triple store.




Figure 5: Direct screenshot of GraphDB Workbench's Visual graph for vnedu:University. It shows the asserted `rdfs:subClassOf` links to higher-education and organization concepts, the external class alignment to `dbo:University`, and the class typing edge. Labels are rendered from the RDF labels stored in the project ontology.




Figure 6: Direct screenshot of Protégé’s Object properties view for the committed ontology. The selected governedBy relation is displayed with bilingual labels, its comment, domain EducationalOrganization, range GoverningBody, and its inverse relation. This is the property-signature view used to verify that the graph has database-like domain–relationship–range discipline; it is not a hand-drawn diagram.





Figure 7: Direct screenshot of GraphDB Workbench using the saved configuration VN-Edu ontology domain-range relations. The configuration expands a class using the ontology’s asserted object-property domains and ranges, so the displayed edges include governedBy, memberOf, ownedBy, ownership, offersProgram, offeredBy, hasMember, trainsMajor and related links. The graph is backed by the local repository vnedu-ontology-d, not generated from a separate illustration.



PublicInstitution and PrivateInstitution are defined by ownership values;
they are orthogonal to the four legal-type classes, not their replacements.

Figure 8: Selected organizational hierarchy. Arrows point from subclass to superclass. The four legal-type classes are no longer globally disjoint: an enduring institution can change type. This is a selected view of 39 classes.





Figure 9: Selected domain–property–range signatures. These arrows are not mandatory foreign keys. Program providers and majors are no longer functional: multiple values must not force different institutions or subjects to become identical.



A current-boundary crosswalk is not a dated history.
The analogous bornIn chain does not prove that P_n existed at birth.

Figure 10: Institution-location inference under the selected administrative crosswalk. It does not establish historical validity. Birth-area mapping now uses a separate birthAreaInCurrentCrosswalk predicate instead of adding historical bornIn statements.

6.3 Implemented corrections and evidence limits

The design is standards-based and suitable for an educational knowledge-graph prototype, but consistency is not a certificate of factual accuracy. Version 2.2 makes the following changes without reminting existing entity identifiers.

Table 3: Ontology 2.2 corrections and their limits.

Concern	Implemented correction	Evidence boundary
Supervision versus inclusion	Removed parent-organization alignments from governance, subordination, ownership and membership. The branch alignment remains because a branch is a structural subdivision.	Supervision or investment does not automatically establish organizational parenthood.
Direct versus indirect governance	Added <code>directlyGovernedBy</code> and <code>reportedGovernedBy</code> under the broad <code>governedBy</code> ; only the broad relation propagates through superiors.	Legacy source relations migrate to reported, not unsupported direct, governance.
Leadership versus employment	Removed leader-to-employee and leader-to-worksFor alignments; added role observations.	A council chair need not be an employee.
Institutional type changes	Removed global disjointness between the four legal-type classes; added classification observations.	The class projection is not a dated authoritative legal register.
Education versus alumni	Recorded education becomes <code>educatedAt</code> , with the neutral inferred class <code>EducationParticipant</code> .	Attendance does not entail graduation. Alumni terms remain for explicit evidence.
Joint programs	Removed functionality from <code>offeredBy</code> and <code>ofMajor</code> ; SHACL requires at least one provider and major.	Curated program identifiers still do not distinguish all curriculum variants.
Birth geography	Successor and containment chains produce <code>birthAreaInCurrentCrosswalk</code> , not additional <code>bornIn</code> facts.	A current province is not asserted to have existed at birth.
Version publication	Archived 2.1 and 2.2 Turtle; generated HTML, Turtle and JSON-LD version documents and dynamic routes.	The source and Pages workflow provide the route; stable public uptime still depends on deployment.

6.4 Standards and conformance assessment

The audited ontology follows the RDF 1.1/Turtle and OWL 2 RDF serialisation conventions and uses explicit OWL classes, object properties, datatype properties, inverse properties, selected functional properties, disjointness, property chains and defined classes. The current source contains 39 local classes, 38 local object properties, 27 local datatype properties and 862 ontology triples. The visualisation above is generated from those current source counts, so it can be regenerated after an ontology change rather than becoming a stale hand-drawn figure.

The design is suitable for a standards-based academic prototype, but it is not an automatic “five-star” certification. The following claims are deliberately separated:

Table 4: Standards assessment and remaining boundaries.


Area	Evidence in this release	Honest qualification
RDF and URI design	Stable HTTPS namespace, version IRI, Turtle/JSON-LD/N-Triples publication, content negotiation checks and licence metadata.	Deployment and long-term availability at the canonical origin remain operational responsibilities.
OWL 2 modelling	The ontology parses as RDF and passes the OWL API OWL2RLProfile gate with zero violations; inverse relations, chains, restrictions, equivalent classes and selected disjointness remain explicit.	This certifies profile membership of the local ontology, not complete imported-vocabulary semantics or factual correctness.
Constraint validation	SHACL shapes validate required fields, datatypes, temporal intervals and observation structure; the current suite reports zero violations under its policy.	Warnings reflect incomplete source coverage; SHACL validation is not proof of factual truth.
Provenance and time	PROV-O-style source links and reified evidence observations preserve snapshot context and optional temporal qualifiers.	The current Silver snapshot still lacks claim-level dates for many facts; an undated observation is not a current fact.
Vocabulary reuse	Conservative links to Schema.org, FOAF, DBpedia, SKOS, PROV-O and external identifiers.	External ontologies are not fully imported and reasoned over; alignments must remain semantically narrow.
Linked Data stars	Machine-readable publication, stable links, external links, licence and open formats are implemented.	The five-star scheme is a publication maturity heuristic, not a formal conformance test, and does not certify accuracy, security or production performance.

6.5 Qualified observations and temporal interpretation

The transformer creates 3,153 `SourceObservation` records, including leadership, measurement and classification specializations. Each reifies a triple through `rdf:subject`, `rdf:predicate` and `rdf:object`. Its SHA-256-derived URI includes the assertion, snapshot digest, source context and qualifiers. Identical evidence preserves the observation URI; changed evidence creates a different record. Entity URIs remain stable.

The implementation supports `validFrom`, `validThrough`, `referenceYear` and measurement units, validates calendar dates and interval order, and retains normalized SHA-256 digests of Silver inputs in `data/reports/observation-provenance.json`. Copied source references provide *entity-level context*, not proof that every assertion occurs in every referenced source. The snapshot digest is not an upstream publication date.

The present Silver snapshot does not retain sufficient temporal qualifiers for these measurements and appointments. Generated observations explicitly have unknown temporal validity. The system does not infer today's leadership or a current enrollment total; the institution infobox displays this limitation. Flat triples remain a source-snapshot projection. Historical queries must require relevant date bounds on observations; absence of an end date does not prove an appointment continues.



No supplied dates $\Rightarrow$ unknown validity, not current validity.

Figure 11: Observation model. Dashed fields are optional and must come from evidence. Missing dates mean unknown temporal validity. The diagram is schematic; no dates are fabricated for real officeholders.

SHACL checks observation structure, allowed status, digest syntax, date datatypes and interval ordering. Regression tests also verify that reification alone does not assert the reified triple. This is a conservative migration, not a complete longitudinal collection system. Acquiring actual dates, claim-level provenance, foundation events and authoritative historical classifications remains future work.

6.6 Competency questions and prohibited conclusions

1. A direct body followed by its superior entails broad governance, not a second direct edge.
2. A council chair does not entail employment; an owner does not entail parenthood.
3. Joint providers or majors do not become identical through a program.
4. Recorded education does not entail the alumni class.
5. A current geographic crosswalk does not create a historical birthplace assertion.
6. An undated observation does not answer a date-valid question merely because it is in the latest export.

Rule-engine execution, formal OWL profile membership, domain correctness and empirical coverage remain distinct claims. The formal OWL 2 RL gate is now asserted only for the committed release ontology, with its reproducible command and zero-violation evidence recorded in `data/reports/owl2rl-profile.txt`.

6.7 Materialization

The system computes a rule closure over the local ontology and asserted data and retains selected new instance-level statements. External identity links are joined into the serving graph after this step. This avoids equality-driven copying to remote identifiers, but means the published union is not a complete closure including all `owl:sameAs` effects. The reported inferred-triple count describes the retained subset, not the full internal closure.

Representative rules are:

$$\begin{aligned} \text{locatedIn} \circ \text{partOf} &\sqsubseteq \text{locatedIn}, \\ \text{locatedIn} \circ \text{mergedInto} &\sqsubseteq \text{locatedIn}, \\ \text{governedBy} \circ \text{subordinateTo} &\sqsubseteq \text{governedBy}, \\ \text{offersProgram} \circ \text{ofMajor} &\sqsubseteq \text{trainsMajor}. \end{aligned}$$

For example, an institution assigned to a former province can also be related to the successor province and its larger region. Inverse properties expose the same relationship in the opposite direction without manually storing both directions in the asserted graph.

The model uses the owlrl implementation of OWL 2 RL/RDF rules. A separate OWL API OWL2RLProfile gate now checks the committed Turtle ontology; release 2.2 produced zero violations. Calendar years are represented as `xsd:integer`, and calendar dates as `xsd:dateTime` at midnight when the source supplies date precision only. This certifies the ontology’s syntactic profile membership, not factual correctness, reasoner completeness over unimported external ontologies, or production performance.

6.8 Geographic and temporal limitations

The `birthAreaInCurrentCrosswalk` chains are geographic containment crosswalks, separated from `bornIn` in version 2.2. Relating a birthplace to a province created in 2025 does not mean that province existed when the person was born. Similarly, historical institutions should not automatically be counted as currently operating just because their location can be mapped to a current province. Queries for current institutions should explicitly consider dissolution evidence and its incompleteness.

Version 2.2 adds qualified observations for roles, measurements and classifications, but the available source snapshot still lacks sufficient temporal qualifiers and historical geographic versions. Absence of an end date is not proof of a current appointment. Birth-date collection currently loses source precision, and suppressing January 1 dates is only a heuristic: it can suppress real birthdays and cannot establish the precision of all other dates. These source-quality limitations remain documented work, not conclusions resolved by successful reasoning.

7 Linking and publication

7.1 External relations

The link graph contains 3,558 triples. Wikidata, ROR, GeoNames, and selected DBpedia links use `owl:sameAs`; subject mappings use `skos:closeMatch`; Wikipedia article links use `foaf:isPrimaryTopicOf`. The last of these relates an entity to a document and must not be counted as an identity assertion. Repeated attempts to add the same link no longer inflate linkset counters.

Shared identifiers are useful identity evidence, but errors and historical scope mismatches can occur in the source mappings. Some organizational links are discovered through label matching with contextual filters. These deserve separate manual review. The audit does not certify every identity link or every administrative boundary match.

7.2 Representations and deployment

The static publisher generates HTML pages, embedded JSON-LD, and RDF files. The dynamic server provides query handling and RDF content negotiation. The canonical namespace currently belongs to

GitHub Pages; the dynamic service has another origin. Content negotiation at the dynamic service does not establish that the canonical static URI performs the same negotiation.

The implementation supports practices associated with five-star Linked Open Data. Section 9 assesses each level and reports a small, dated HTTP sample; it does not reuse the earlier report's live availability measurements. The repository declares data licensing and retains media metadata; this audit is not a legal verification of the license of every upstream asset.

The publisher now checks its output path before recursive replacement. It refuses the project root, external paths, symbolic links, and existing directories without generated-site markers. This prevents a misconfigured output environment variable from deleting source data or the project itself.

8 URI publication and Apache Jena Fuseki

Fuseki is a SPARQL server, not a URI minting service [9]. The transformer assigns canonical identifiers before data enters any store. The same university IRI remains in RDF whether queried by RDFLib, Oxigraph or Fuseki.

8.1 Identity, representation and storage

```
# RDF subject identifying a university:
https://pham-ng.github.io/Vietnam-University-Knowledge-Graph-ver2/resource/university/dai-hoc-j
↪ bach-khoa-ha-noi
# Local query endpoint:
http://localhost:3030/vnedu/sparql
# Graph Store Protocol target:
http://localhost:3030/vnedu/data?default
```

Uploading RDF does not make its canonical IRIs dereferenceable automatically. The publisher independently creates HTML, Turtle and JSON-LD descriptions. Ontology terms retain hash IRIs; entity and observation identifiers use slash paths. Immutable 2.1 and 2.2 version artifacts are now generated under `ontology/`. Dynamic version routes support RDF content negotiation and return 404 for unknown versions.

GitHub Pages remains a separate HTTP origin from Flask. Its static representation links and embedded JSON-LD do not implement Flask's negotiation, nor the entity-to-document 303 pattern discussed in *Cool URIs* [6]. Adding Fuseki cannot fix those origin-level behaviors. Version artifacts are implemented and tested locally; their production deployment requires promotion of the review branch.

8.2 Safe loading and reproducible runtime

The revised loader requires a successfully validated release before network writes. It targets an existing explicitly named dataset rather than creating one implicitly. An empty graph can be loaded directly. Replacement of a nonempty default graph requires both `--replace` and a new `--backup` path; the backup is created exclusively, so an older backup cannot be silently overwritten. HTTP redirects and error responses are rejected. After PUT, the loader retrieves the entire default graph and checks RDF graph isomorphism, including blank nodes, rather than relying on a triple count.

Graph Store PUT replaces the selected graph [11]; named graphs are untouched by this loader. Credentials are read from environment variables, not command-line passwords. A failed comparison preserves the backup for operator recovery; it is not a distributed transaction or automatic rollback. Operators must stop concurrent writers during a release: a backup followed by PUT does not provide compare-and-swap concurrency control.




Figure 12: Publication and storage paths. The optional Fuseki branch has verified Linux and Windows acceptance runs; Render still explicitly selects local RDFLib. Storage does not replace the canonical identifiers.

The runtime setup pins Apache Jena Fuseki 6.2.0 and verifies its published SHA-512 digest. The local launcher requires Java 21, validates the release, binds to loopback, and enables only query access with a configured timeout. Its file mode reloads the graph at startup; the separate acceptance harness explicitly tests persistent TDB2. The optional legacy Docker recipe now requires an operator-selected image and password, and binds port 3030 only on loopback. That Docker configuration was not the tested runtime.

8.3 Production-like security and performance acceptance

The additional `audit/production_readiness.py` test starts the same Waitress WSGI entry point used by `app/server.py --prod` with the read-only local backend. Under normal bounded concurrency (eight requests, two query workers), all eight health requests returned HTTP 200; measured latency was 4,213 ms at p50, 4,873 ms at p95, and 4,873 ms maximum on the Windows review machine. An intentional eight-request overload returned two HTTP 200 responses and six HTTP 503 responses, demonstrating bounded back-pressure rather than unbounded queue growth. Security probes rejected a path-injection attempt (404), an unapproved SSRF SERVICE target (400), and an oversized query (413); the response included `nosniff`, `SAMEORIGIN`, and `strict-referrer-policy` headers.

These numbers are a reproducible loopback acceptance sample, not a capacity SLO, penetration test, TLS/reverse-proxy audit, cloud IAM review, or production security certification. The process-per-query RDFLib backend deliberately trades throughput for isolation; a real deployment still needs external rate limiting, resource quotas, TLS termination, monitoring, authentication policy, and a load test against its selected Fuseki/WSGI topology. The full machine-readable evidence is `data/reports/production-readiness.json`.

8.4 Executed end-to-end acceptance evidence

The primary reproducible test ran on Linux GitHub Actions with Fuseki 6.2.0 and Temurin Java 21.0.12.1, on 7 October 2026 at 01:50 UTC. A second run on the Windows desktop used the same Fuseki JAR, Java 21 and the Windows selector/socket workaround at 07:22 UTC; the launcher smoke test passed at 07:31 UTC. The tested code revision before these launcher edits was `ab0033d`; the current branch records the subsequent launcher fix. The input hash, JAR hash, runtime details and assertions are archived in `data/reports/fuseki-e2e-linux.json` and the Windows run in `data/reports/fuseki-e2e-windows.json`; the corresponding Linux CI run completed successfully.

Table 5: Measured Fuseki/TDB2 acceptance results.

Check	Observed result
Load complete validated serving graph	73,690 triples; graph-isomorphic to local RDF.
Attempt unapproved replacement	Refused before PUT.
Approved replacement with backup	Backup graph-isomorphic to previous contents.
Query parity with RDFLib	Count: 1 row; class totals: 68 rows; programs: 78 rows; explicitly dated leadership: 0 rows.
Stop and restart with the same TDB2 directory	Entire default graph remains isomorphic.
Restart without update permission	SPARQL Update and Graph Store PUT both return HTTP 405.
Read after rejected writes	Graph unchanged and isomorphic.

The test uses a newly created disposable directory and ephemeral loopback port; no existing user dataset is overwritten. Zero dated-leadership rows is an honest reflection of missing date evidence, not a failed lookup to be filled with invented dates.

On Windows, the first attempt exposed a JDK 21 selector/socket initialization failure. The repaired launcher now sets the selector provider and short Unix-domain socket directory through inherited Java options (avoiding PowerShell 5.1 argument concatenation); the same disposable TDB2 acceptance test then passed all checks, and a direct query-only launcher smoke test returned HTTP 200. This proves local Windows operability, not production deployment. Configured timeout is not proof of cancellation under load. Authentication, SERVICE egress, concurrent writers, memory limits, disaster recovery and production TLS require separate acceptance tests. The verified read-only mode is access restriction, not a full security certification.

9 Five-star Linked Open Data assessment

The five-star scheme evaluates how data is published and connected, not the overall quality of a program or the correctness of every ontology axiom [7]. There is no Fuseki-specific requirement: RDFLib or static RDF publication can support the relevant mechanisms. Table 6 assesses the project’s data publication rather than assigning an unsupported software-quality score.

Table 6: Evidence against the cumulative five-star publication scheme.

Level	Criterion, summarized	Project evidence and qualification
1	Publicly accessible, openly licensed	Public sample accessible; dataset metadata declares CC BY-SA 4.0.
2	Structured, machine-readable	Upstream media retain their own licenses. JSON records and parsed RDF graphs, not only screen images or PDFs.
3	Non-proprietary format	JSON, Turtle and JSON-LD artifacts.
4	Web identifiers and open Semantic Web standards	Stable-registry HTTP IRIs, RDF descriptions and query interfaces; sampled entity documents parse. Canonical negotiation remains static; version artifacts are now implemented on the review branch.
5	Connections to other datasets	Corrected local link graph contains 2,309 owl:sameAs, 61 skos:closeMatch and 1,188 document-topic links. Existence is measured; every identity assertion is not independently verified.

Verdict. The implementation and sampled public RDF support the basic five-star publication pattern, subject to the declared licensing scope. This is not an independent certification, proof of factual accuracy, or proof that the corrected branch is deployed. Deployment, temporal-evidence and operations limits should be reported alongside the claim, not hidden by a five-star badge. Failing a single endpoint probe does not automatically reduce an otherwise accessible RDF dataset to a lower star level.

9.1 Dated HTTP evidence

`audit/check_publication.py` performs seven read-only requests with bounded timeouts; the result is saved in `data/reports/publication-check.json`. This historical, pre-migration sample was recorded on 7 October 2026 at 00:10 ICT (UTC+7). The institution tested was `resource/university/dai-hoc-bach-khoa-ha-noi`; this is not a random or exhaustive sample.

Table 7: Observed HTTP behavior; deployed content may predate the corrected branch.

Request	Observation
Canonical entity, Accept HTML	HTTP 200, HTML, embedded JSON-LD and alternate RDF links.
Canonical entity, Accept Turtle	HTTP 200, still HTML; no HTTP redirect.
Entity's .ttl document	HTTP 200, Turtle, 204 triples; graph describes the entity IRI.
Entity's .jsonld document	HTTP 200, JSON-LD, 204 triples; graph describes the entity IRI.
Declared ontology/2.2 version	Source build emits immutable HTML, Turtle and JSON-LD; live Pages availability awaits deployment of this corrected branch.
Render /healthz	Read timeout after 15 seconds; health and deployed commit were not established.
Local Fuseki /\$/ping	Connection refused on port 3030; no reachable service at the configured address.

A Render timeout can reflect a cold start or a network/service problem; it is not proof of permanent downtime. The probe did not test SPARQL federation, persistence, admin credentials, or the full set of resource identifiers. The existing public deployment remains distinct from the audited branch, whose Linux CI and local tests validate a different scope.

The media distinction also matters. For example, the shipped HUST logo metadata records “fair use” while its campus photograph records CC0. Describing or linking such an image does not relicense the image under the dataset license. The revised publisher shows a source link, not an embedded image, for unknown, fair-use or unrecognized licenses. A data publication claim must not be expanded into a blanket claim that every remotely displayed asset is openly licensed.

10 Query execution and security review

The original SPARQL policy searched for textual patterns containing whitespace after `SERVICE` and `FROM`. Valid forms without that whitespace bypassed the filter. The audit reproduced this locally using loopback targets without issuing an outbound request. The revised policy inspects parsed SPARQL operations and permits only two exact HTTPS service endpoints. It rejects dataset-loading clauses, variable service targets, alternate ports, credentials, and arbitrary paths. Security keywords inside string literals are not treated as operations. The local outbound HTTP handler also rejects redirects.

The previous timeout waited only for `Graph.query` to return. `RDFLib` can evaluate `SELECT` bindings lazily, so result evaluation could continue beyond the limit. Cancelling a future also cannot stop an already running Python thread. The revised local backend evaluates and serializes each query in a disposable process, terminates it on timeout, and limits concurrent workers. An expensive Cartesian-product regression test checks that capacity is released and a subsequent small query succeeds.

The local backend limits serialized responses to 8 MiB. This does not impose a hard memory limit during intermediate evaluation, and parsing the public request still occurs in the web process. A hostile-workload production deployment therefore needs process or container memory limits, request-rate controls, and appropriate upstream isolation. The process-per-query approach also reloads the graph and has a throughput cost. Earlier in-memory latency measurements do not describe this implementation.

The Fuseki proxy supplies an execution timeout parameter and does not follow proxy redirects. Actual query cancellation and remote SERVICE egress remain responsibilities of the configured Fuseki instance. The separate Linux acceptance run verifies load/query/persistence/read-only behavior; it does not verify production federation controls. Raw POST requests with application/sparql-query are now accepted, alongside form-encoded query requests.

11 Evaluation and results

11.1 Release measurements

All values in Table 8 are generated by `audit/release_metrics.py` from the validated local release. That script also checks that the serving graph is isomorphic to the union of its declared components. It does not request data from a remote endpoint.

Table 8: Measured size of the corrected local release.

Quantity	Count
Educational-organization records	300
Classified higher-education institutions	271
Person records	1,702
Institution-leader records	221
People with recorded education relations	1,481
Current / former province-level units	34 / 29
Illustrative programs / majors	78 / 37
Asserted data triples	47,291
Retained inferred triples	21,825
Link-graph triples	3,558
Ontology / dataset-metadata triples	862 / 96
Distinct triples in serving union	73,628
Local resource IRIs appearing as subjects	5,401

In the corrected snapshot, selected SHACL checks report 0 violations, 34 warnings, and 1 informational result. Warnings include missing foundation years and institution locations. Warning-tolerant conformance means the configured release policy permits these omissions; it does not mean all records are complete. The consistency diagnostics report no detected local violations after the fixes.

Field coverage uses all 300 institution records as its denominator: 278 have a selected foundation year, 288 a province reference, 256 an ownership value, 237 a website, 175 coordinates, 60 an admission code, and 28 a student-count value. These are presence measurements, not estimates of truth, timeliness, or national coverage. In particular, 271 higher-education classifications and 300 educational-organization records are different denominators and should not be interchanged.

11.2 Link experiment: agreement, not independent accuracy

The available experiment uses 86 institutions with English labels and identifier-derived DBpedia reference targets. A local token-Jaccard matcher selects the highest-scoring target from that closed reference candidate set. Table 9 reports agreement with the reference frozen from revision 58c65f0. This baseline is not Silk; the separate actual Silk execution is reported below.

Let M be the set of predictions, R the reference pairs, and $TP = |M \cap R|$. The reported precision is $TP/|M|$, recall is $TP/|R|$, and F1 is their harmonic mean. Calling R independent ground truth would

Table 9: Token-Jaccard agreement with the identifier-derived reference.

Threshold	Predictions	Disagreements	Ref. precision	Ref. recall	Ref. F1
0.2	85	13	84.7%	83.7%	0.842
0.4	80	8	90.0%	83.7%	0.867
0.6	76	6	92.1%	81.4%	0.864
0.8	68	4	94.1%	74.4%	0.831
1.0	58	1	98.3%	66.3%	0.792

be incorrect because the identifier linker generated it. Scoring that same linker against R yields perfect agreement by construction. The previous claim of 100% linking accuracy is therefore withdrawn.

The selected sample excludes unlinked institutions and targets outside the reference set. These restrictions affect task difficulty and coverage. The threshold with the highest F1 on this sample is a descriptive observation, not a validated threshold chosen on a separate development set. A defensible accuracy study needs independently annotated links and non-links, dated evidence, an adjudication procedure, and a held-out test set.

11.3 Executed Silk Framework experiment

A separate experiment now executes actual Silk Framework 3.6.0, rather than referring to the Python Jaccard baseline as Silk. The official single-machine entry point was compiled with Scala 2.12.10 against released workbench libraries and run on Java 8. Pinned checksums are in `audit/setup_runtimes.ps1`; the runner, XML rules, input N-Triples, candidate outputs and result JSON are included in the repository. The reproducible execution is scheduled in `.github/workflows/silk-e2e.yml` on Ubuntu with Java 8. On the reviewed Windows desktop, Java 8.351 and 8.481 both timed out before producing candidates; this is recorded as an environment limitation, not converted into a success claim.

The experiment freezes the 86-source/86-target sample from revision 58c65f0, giving 7,396 possible pairs. Both methods use the same lowercased English labels. Silk uses edit-distance thresholds of 0, 3 and 8 with a top-one filter, while the Python baseline uses token-Jaccard similarity. Their threshold values are not numerically comparable. The XML filter is inside `LinkageRule`, as required by the 3.6.0 parser; assertions verify at most one prediction per source and the experiment-only predicate. Older example configurations are not assumed equivalent.

Table 10: Actual Silk versus the fixed-sample Python baseline: reference agreement only.

Rule	Predicted	Matches	Ref. precision	Ref. recall	Ref. F1
Silk: distance 0	57	57	100.0%	66.3%	0.797
Silk: distance 3	64	63	98.4%	73.3%	0.840
Silk: distance 8	75	70	93.3%	81.4%	0.870
Jaccard: similarity 0.4	80	72	90.0%	83.7%	0.867
Jaccard: similarity 0.6	76	70	92.1%	81.4%	0.864
Jaccard: similarity 0.8	68	64	94.1%	74.4%	0.831

These results do not estimate independent identity accuracy. Some historical reference identities are rejected by the current stricter linker; they remain in the frozen experimental reference, not in accepted release identities. The sample excludes open-world negatives, lacks independent adjudication and has no held-out tuning set. Silk’s tie handling is engine-defined; Python uses descending target URI for ties. Small differences do not establish general superiority.

No prediction is promoted to `owl:sameAs`. Silk emits a candidate-only URN predicate and its files are outside the serving pipeline. It is a research extension, not a normal publication dependency. Neo4j is not

added: a second graph model and synchronization path would supply neither missing time evidence nor another five-star requirement. The implementation follows the official Silk source and link-specification documentation [8].

11.4 Regression evidence

The audit adds counterexamples for compact SPARQL syntax, exact service endpoint policy, keywords inside literals, raw query POSTs, cancellation and worker recovery, datatype diagnostics, literal ranges, publication gates, stale release records, persistent URIs, same-name identity collisions, offline cache misses, impossible dates, and dangerous publisher paths. Eight initial counterexamples failed before correction and passed after the relevant fixes. The final test-run result and command are recorded in the companion audit notes rather than a manually maintained badge.

The full pipeline is reconstructed from the shipped cache, and both static-site variants are built before the server integration suite. A successful run establishes that the tested release follows the intended local contracts. It does not constitute a penetration test, an independent factual annotation study, or a load test of the public deployment.

11.5 Treatment of earlier comparisons

The earlier report's 21-fact comparison with a predecessor repository is a small, hand-selected demonstration. Its benchmark script relies on a machine-specific path and permissive answer matching, and the underlying cases do not have a complete independent annotation protocol. Its scores are not used here as an accuracy estimate. Claims about 349 predecessor contradictions, graph-density superiority, live link availability, and public-endpoint latency are likewise not repeated as newly verified findings.

Graph density is especially sensitive to which nodes and predicates are counted. Adding external nodes or inferred inverse edges changes the measurement without necessarily adding useful or correct information. A future comparison should define a common node universe, edge projection, deduplication policy, and frozen input revisions before reporting degree statistics.

12 Threats to validity and remaining work

Construct validity. Logical consistency, SHACL conformance, field presence, and link agreement measure different things. None alone measures real-world correctness. Program examples and name-based legal-type classification further restrict what answers can be interpreted as authoritative.

Internal validity. Source precedence can select a plausible but incorrect value. Earliest-year selection conflates events unless the event meaning is modeled explicitly. Some unqualified dates and organizational relationships need source-level adjudication. The revised person split prevents unsupported merges, but may also separate records belonging to the same real person; this is an explicit uncertainty trade-off.

External validity. Coverage favors institutions and people documented in the selected community sources. The 86-link experiment omits open-world candidates and cannot be generalized to all Vietnamese institutions. National completeness requires comparison with a dated authoritative institutional register.

Operational validity. Local tests do not prove public availability, deployed revision identity, traffic capacity, or Fuseki egress behavior. Memory limits and endpoint rate controls remain deployment concerns. Snapshot timestamps record reconstruction or collection events as implemented; they must not automatically be interpreted as the dates on which the underlying facts became true.

The next substantive data work is acquiring real temporal qualifiers for the new observations, event-specific foundation histories, and independently verified assertion-level provenance. A subsequent research evaluation should use a preregistered sampling method, an independently reviewed reference set, and a clear distinction between discovery coverage, linking coverage, and correctness. Those additions require evidence not present in the current repository and are not supplied by inventing dates, labels, or outcomes.

13 Conclusion

VN-Edu LOD demonstrates a working integration and publication pipeline for a Vietnamese education knowledge graph. The audited implementation strengthens the boundary between untrusted queries and execution, between candidate data and publication, and between source labels and entity identity. Its measured result is a locally validated release with 73,628 distinct serving triples and explicit limitations on factual, temporal, and operational interpretation. The available evidence supports the system as an educational Linked Data artifact; it does not establish perfect accuracy or complete coverage of Vietnamese higher education.

A Reproducing the release and report

Use Python 3.12 and the pinned dependencies. After installation, the default pipeline uses the shipped cache and fails on missing cached responses. Run from the repository root:

```
python -m pip install -r requirements.txt
python run_all.py --no-install
python audit/release_metrics.py
```

Build the dynamic-server UI before running the complete integration suite. In PowerShell:

```
$env:VNEDU_SITE_DIR = 'site_server'
$env:VNEDU_SITE_ROOT = '/'
python scripts/step7_publish.py
python -m pytest tests -q
```

Unset the two UI environment variables before building the GitHub Pages variant. The default pipeline already builds that variant. Existing output directories must have the publisher's generated-site markers.

From docs/report, compile the English edition with an installed TeX distribution:

```
latexmk -lualatex -interaction=nonstopmode -halt-on-error main-en.tex
```

The PDF uses Libertinus and Noto Sans Mono. The numeric macros come from `tables/audit-metrics.tex`; regenerate them after data or ontology changes. `data/reports/audit-metrics.json` contains the corresponding machine-readable values. The original Vietnamese source and PDF remain available as historical project artifacts.

B A competency-query example

This query counts classified institutions by current province using the retained geographic inference. It excludes institutions explicitly marked as defunct; the absence of that marker still does not prove current operation.

```
PREFIX vnedu: <https://pham-ng.github.io/Vietnam-University-Knowledge-Graph-ver2/ontology#>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
SELECT ?province (COUNT(DISTINCT ?institution) AS ?count)
WHERE {
  ?institution a vnedu:HigherEducationInstitution ;
               vnedu:locatedIn ?p .
  ?p a vnedu:Province ; rdfs:label ?province .
  FILTER(LANG(?province) = "vi")
  FILTER NOT EXISTS { ?institution a vnedu:DefunctInstitution }
}
GROUP BY ?province
ORDER BY DESC(?count)
```

The example illustrates an essential reporting distinction: a query can execute correctly against the model while its answer remains dependent on the completeness and meaning of the underlying evidence.

References

- [1] R. Cyganiak, D. Wood, and M. Lanthaler, eds. *RDF 1.1 Concepts and Abstract Syntax*. W3C Recommendation, 2014. <https://www.w3.org/TR/rdf11-concepts/>.
- [2] B. Motik et al., eds. *OWL 2 Web Ontology Language Profiles*, second edition. W3C Recommendation, 2012. <https://www.w3.org/TR/owl2-profiles/>.
- [3] H. Knublauch and D. Kontokostas, eds. *Shapes Constraint Language (SHACL)*. W3C Recommendation, 2017. <https://www.w3.org/TR/shacl/>.
- [4] S. Harris and A. Seaborne, eds. *SPARQL 1.1 Query Language*. W3C Recommendation, 2013. <https://www.w3.org/TR/sparql11-query/>.
- [5] T. Lebo, S. Sahoo, and D. McGuinness, eds. *PROV-O: The PROV Ontology*. W3C Recommendation, 2013. <https://www.w3.org/TR/prov-o/>.
- [6] L. Sauermann and R. Cyganiak, eds. *Cool URIs for the Semantic Web*. W3C Interest Group Note, 2008. <https://www.w3.org/TR/cooluris/>.
- [7] T. Berners-Lee. *Linked Data*, including the five-star scheme. Design Issues, 2006, scheme added 2010. <https://www.w3.org/DesignIssues/LinkedData.html>.
- [8] Silk Framework contributors. *Silk Framework 3.6.0: Link Specification Language and implementation*. <https://github.com/silk-framework/silk/tree/v3.6.0>. Accessed 7 October 2026.
- [9] Apache Jena. *Apache Jena Fuseki*. Project documentation, accessed 7 October 2026. <https://jena.apache.org/documentation/fuseki2/>.
- [10] Apache Jena. *Fuseki HTTP Administration Protocol*. <https://jena.apache.org/documentation/fuseki2/fuseki-server-protocol.html>.
- [11] C. Ogbuji, ed. *SPARQL 1.1 Graph Store HTTP Protocol*. W3C Recommendation, 2013. <https://www.w3.org/TR/sparql11-http-rdf-update/>.
- [12] Apache Jena. *Security in Fuseki2*. <https://jena.apache.org/documentation/fuseki2/fuseki-security.html>.
- [13] Schema.org. *parentOrganization; employee*. Vocabulary definitions, accessed 7 October 2026. <https://schema.org/parentOrganization>; <https://schema.org/employee>.